

```
%%writefile evaluate_fitness_batch.py
import numpy as np
import pandas as pd
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Load dataset
df = pd.read_csv("Crop_recommendation.csv")
X = df.drop(columns="label") # Ensure all features are included
y = df["label"]

# Read population (batch of chromosomes) from a file
population = np.loadtxt("population_batch.txt", delimiter=",", dtype=int)
print(f"Population Shape: {population.shape}")

# Validate chromosome length
if population.shape[1] != X.shape[1]:
    raise ValueError(f"Chromosome length {population.shape[1]} does not match the number of

# Initialize list to store fitness values
fitness_results = []

# Loop through each chromosome in the population
for chromosome_id, chromosome in enumerate(population, start=1):
    print(f"Chromosome {chromosome_id}: {chromosome}")

    # Get selected features
    selected_features = np.where(chromosome == 1)[0]
    print("Selected Features:", selected_features)

    # Handle case where no features are selected
    if len(selected_features) == 0:
        print(f"Chromosome {chromosome_id}: No features selected. Assigning fitness 0.")
        fitness_results.append(0.0)
        continue

    # Select features based on chromosome
    X_selected = X.iloc[:, selected_features]

    # Train/test split
    X_train, X_test, y_train, y_test = train_test_split(X_selected, y, test_size=0.3, random

    # Train the model
    rf = RandomForestClassifier(n_estimators=50, random_state=42, criterion="entropy")
    rf.fit(X_train, y_train)

    # Predict and calculate accuracy
    y_pred = rf.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
```


```

print(f"Chromosome {chromosome_id} Accuracy:", accuracy)

# Store fitness result
fitness_results.append(accuracy)

# Save all fitness results to a file
np.savetxt("fitness_results_batch.txt", fitness_results, fmt="%.4f")
print("Fitness results saved.")

```

 Overwriting evaluate_fitness_batch.py

```

%%writefile cudamain.cu

#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>
#include <curand_kernel.h>

#define POP_SIZE 200
#define CHROMOSOME_LENGTH 7
#define NUM_GENERATIONS 50
#define MUTATION_RATE 0.01f
#define MAX_STAGNATION 3
#define BLOCK_SIZE 256

typedef struct {
    int id;
    int genes[CHROMOSOME_LENGTH];
    double fitness;
} Individual;

// CUDA kernel for random chromosome generation
__global__ void generate_population_kernel(int* population, curandState* states,
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < pop_size) {
        curandState localState = states[idx];
        for (int i = 0; i < chromosome_length; i++) {
            population[idx * chromosome_length + i] = curand(&localState) % 2;
        }
        states[idx] = localState;
    }
}

// CUDA kernel for mutation
__global__ void mutation_kernel(int* population, curandState* states, float mutat
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < pop_size * chromosome_length) {
        int chr_idx = idx / chromosome_length;
        int gene_idx = idx % chromosome_length;

```

```

        curandState localState = states[chr_idx];
        float r = curand_uniform(&localState);
        if (r < mutation_rate) {
            population[idx] = 1 - population[idx];
        }
        states[chr_idx] = localState;
    }
}

// CUDA kernel for crossover
__global__ void crossover_kernel(int* population, int* new_population, int pop_size,
    int idx = blockIdx.x * blockDim.x + threadIdx.x;

    if (idx < pop_size / 2) {
        int parent1_idx = 2 * idx;
        int parent2_idx = 2 * idx + 1;

        int swap_start = chromosome_length - 3;

        // Pointers to parents and offspring
        int* parent1 = &population[parent1_idx * chromosome_length];
        int* parent2 = &population[parent2_idx * chromosome_length];
        int* offspring1 = &new_population[parent1_idx * chromosome_length];
        int* offspring2 = &new_population[parent2_idx * chromosome_length];

        // Perform three-point crossover
        for (int i = 0; i < chromosome_length; i++) {
            if (i < swap_start) {
                offspring1[i] = parent1[i];
                offspring2[i] = parent2[i];
            } else {
                offspring1[i] = parent2[i];
                offspring2[i] = parent1[i];
            }
        }
    }
}

// Function to initialize random states
__global__ void init_random_states(curandState* states, unsigned long seed) {
    int idx = blockIdx.x * blockDim.x + threadIdx.x;
    if (idx < POP_SIZE) {
        curand_init(seed, idx, 0, &states[idx]);
    }
}

int main() {
    // Allocate device memory
    int *d_population, *d_new_population;
    curandState *d_states;
    cudaMalloc(&d_population, POP_SIZE * CHROMOSOME_LENGTH * sizeof(int));
    cudaMalloc(&d_new_population, POP_SIZE * CHROMOSOME_LENGTH * sizeof(int));

```

```

cudaMalloc(&d_states, POP_SIZE * sizeof(curandState));

// Initialize random states
int blocks = (POP_SIZE + BLOCK_SIZE - 1) / BLOCK_SIZE;
init_random_states<<<blocks, BLOCK_SIZE>>>(d_states, time(NULL));
cudaDeviceSynchronize();

// Generate initial population
generate_population_kernel<<<blocks, BLOCK_SIZE>>>(d_population, d_states, PC
cudaDeviceSynchronize();

// Allocate host memory
Individual* population = (Individual*)malloc(POP_SIZE * sizeof(Individual));
Individual* new_population = (Individual*)malloc(POP_SIZE * sizeof(Individual)
int* h_population = (int*)malloc(POP_SIZE * CHROMOSOME_LENGTH * sizeof(int));
cudaMemcpy(h_population, d_population, POP_SIZE * CHROMOSOME_LENGTH * sizeof(

// Initialize population structs
for (int i = 0; i < POP_SIZE; i++) {
    population[i].id = i + 1;
    memcpy(population[i].genes, &h_population[i * CHROMOSOME_LENGTH], CHROMOS
}

double best_fitness = -1.0;
int stagnation_count = 0;

for (int generation = 0; generation < NUM_GENERATIONS; generation++) {
    printf("\n=== Generation %d ===\n", generation);

    // Write population to file
    FILE* file = fopen("population_batch.txt", "w");
    if (!file) {
        perror("Error opening batch file");
        exit(EXIT_FAILURE);
    }
    for (int i = 0; i < POP_SIZE; i++) {
        for (int j = 0; j < CHROMOSOME_LENGTH; j++) {
            fprintf(file, "%d", population[i].genes[j]);
            if (j < CHROMOSOME_LENGTH - 1) fprintf(file, ",");
        }
        fprintf(file, "\n");
    }
    fclose(file);

    // Execute Python script
    system("python3 evaluate_fitness_batch.py");

    // Read fitness results
    file = fopen("fitness_results_batch.txt", "r");
    if (!file) {
        perror("Error opening fitness results batch file");
    }
}

```

```

        exit(EXIT_FAILURE);
    }
    for (int i = 0; i < POP_SIZE; i++) {
        fscanf(file, "%lf", &population[i].fitness);
    }
    fclose(file);

    // Find the best individual
    double current_best_fitness = -1.0;
    int best_idx = -1;
    for (int i = 0; i < POP_SIZE; i++) {
        if (population[i].fitness > current_best_fitness) {
            current_best_fitness = population[i].fitness;
            best_idx = i;
        }
    }

    printf("Best Fitness: %.4f\n", current_best_fitness);

    // Check for stagnation
    if (current_best_fitness == best_fitness) {
        stagnation_count++;
        printf("No improvement in generation %d. Stagnation count: %d\n", generat
        if (stagnation_count >= MAX_STAGNATION) {
            printf("Stopping early due to stagnation: No improvement for %d gener
            break;
        }
    } else {
        best_fitness = current_best_fitness;
        stagnation_count = 0;
    }

    // Preserve elite
    new_population[0] = population[best_idx];

    // Perform crossover and mutation
    cudaMemcpy(d_population, h_population, POP_SIZE * CHROMOSOME_LENGTH * siz
    crossover_kernel<<<blocks, BLOCK_SIZE>>>(d_population, d_new_population,
    cudaDeviceSynchronize();
    mutation_kernel<<<blocks, BLOCK_SIZE>>>(d_new_population + CHROMOSOME_LEN
    cudaDeviceSynchronize();


    // Copy results back
    cudaMemcpy(h_population, d_new_population, POP_SIZE * CHROMOSOME_LENGTH *
    for (int i = 1; i < POP_SIZE; i++) {
        memcpy(population[i].genes, &h_population[i * CHROMOSOME_LENGTH], CHF
    }
    population[0] = new_population[0];
}

// Cleanup


```

```
    cudaFree(d_population);
    cudaFree(d_new_population);
    cudaFree(d_states);
    free(h_population);
    free(population);
    free(new_population);

    return 0;
}
```

 Overwriting cudamain.cu

```
!nvcc -o cudamain cudamain.cu
```

 **cudamain.cu(37): warning** #177-D: variable "**gene_idx**" was declared but never referenced

```
    int gene_idx = idx % chromosome_length;
        ^
```

Remark: The warnings can be suppressed with "-diag-suppress <warning-number>"

```
import time

t = time.time()
!./cudamain
```



```


Chromosome 174: [1 1 1 0 1 0 1]
Selected Features: [0 1 2 4 6]
Chromosome 174 Accuracy: 0.9878787878787879
Chromosome 175: [0 0 1 0 0 1 0]
Selected Features: [2 5]
Chromosome 175 Accuracy: 0.4681818181818182
Chromosome 176: [1 1 1 0 1 1 0]
Selected Features: [0 1 2 4 5]
Chromosome 176 Accuracy: 0.9409090909090909
Chromosome 177: [1 0 1 1 0 1 1]
Selected Features: [0 2 3 5 6]
Chromosome 177 Accuracy: 0.9575757575757575
Chromosome 178: [1 0 1 1 1 0 0]
Selected Features: [0 2 3 4]
Chromosome 178 Accuracy: 0.9166666666666666
Chromosome 179: [1 0 1 0 0 0 1]
Selected Features: [0 2 6]
Chromosome 179 Accuracy: 0.8424242424242424
Chromosome 180: [0 0 1 1 1 1 1]
Selected Features: [2 3 4 5 6]
Chromosome 180 Accuracy: 0.9742424242424242
Chromosome 181: [1 1 1 1 1 0 0]
Selected Features: [0 1 2 3 4]
Chromosome 181 Accuracy: 0.9560606060606060
Chromosome 182: [1 1 1 0 1 0 0]
Selected Features: [0 1 2 4]
Chromosome 182 Accuracy: 0.9363636363636364
Chromosome 183: [0 1 0 0 1 0 0]
Selected Features: [1 4]
Chromosome 183 Accuracy: 0.6484848484848484
Chromosome 184: [0 1 0 1 0 0 1]
Selected Features: [1 3 6]
Chromosome 184 Accuracy: 0.8545454545454545
Chromosome 185: [0 0 1 1 1 1 1]

```

```

delta_t = time.time()-t
print("Time taken for total execution of parallel code:", delta_t)
parallel_time = delta_t

```

 Time taken for total execution of parallel code: 195.64607548713684

```
!nvprof --print-summary ./cudamain
```



```

Selected Features: [0 1 2 7]
Chromosome 194 Accuracy: 0.9363636363636364
Chromosome 195: [0 0 0 0 1 1 1]
Selected Features: [4 5 6]
Chromosome 195 Accuracy: 0.8621212121212121
Chromosome 196: [0 0 1 1 1 0 1]
Selected Features: [2 3 4 6]
Chromosome 196 Accuracy: 0.9742424242424242
Chromosome 197: [0 0 0 1 0 1 1]
Selected Features: [3 5 6]
Chromosome 197 Accuracy: 0.75
Chromosome 198: [1 1 0 1 0 0 0]
Selected Features: [0 1 3]
Chromosome 198 Accuracy: 0.7090909090909091
Chromosome 199: [1 0 1 1 0 1 1]
Selected Features: [0 2 3 5 6]
Chromosome 199 Accuracy: 0.9575757575757575
Chromosome 200: [1 0 0 0 0 1 0]
Selected Features: [0 5]
Chromosome 200 Accuracy: 0.296969696969697
Fitness results saved.
Best Fitness: 0.9924
No improvement in generation 3. Stagnation count: 3
Stopping early due to stagnation: No improvement for 3 generations.
==19859== Profiling application: ./cudamain
==19859== Profiling result:

```

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	80.25%	200.26us	1	200.26us	200.26us	200.26us	init_ran
	7.86%	19.616us	3	6.5380us	6.4320us	6.7200us	crossove
	4.07%	10.144us	3	3.3810us	3.3280us	3.4560us	mutation
	3.44%	8.5760us	4	2.1440us	1.9840us	2.3680us	[CUDA me
	3.04%	7.5840us	1	7.5840us	7.5840us	7.5840us	generate
	1.35%	3.3600us	3	1.1200us	1.1200us	1.1200us	[CUDA me
API calls:	96.63%	176.87ms	3	58.956ms	3.1930us	176.86ms	cudaMall
	2.74%	5.0171ms	8	627.13us	8.6540us	4.8334ms	cudaLaun
	0.18%	329.04us	3	109.68us	8.6130us	225.35us	cudaFree
	0.18%	328.66us	7	46.950us	24.331us	77.251us	cudaMemc
	0.15%	282.94us	8	35.367us	6.4690us	219.88us	cudaDevi
	0.09%	158.86us	114	1.3930us	170ns	58.221us	cuDevice
	0.01%	26.368us	1	26.368us	26.368us	26.368us	cuDevice
	0.01%	12.713us	1	12.713us	12.713us	12.713us	cuDevice
	0.00%	5.5500us	1	5.5500us	5.5500us	5.5500us	cuDevice
	0.00%	1.6420us	3	547ns	180ns	1.1050us	cuDevice
	0.00%	967ns	2	483ns	215ns	752ns	cuDevice
	0.00%	654ns	1	654ns	654ns	654ns	cuModule
	0.00%	265ns	1	265ns	265ns	265ns	cuDevice

